

CS726 Programming Assignment – 4 Report

Saksham Rathi (22B1003)

Sharvaneer Sonawane (22B0943)

Deeksha Dhiwakar (22B0988)

Department of Computer Science,
Indian Institute of Technology Bombay

Task 0: Environment Setup and Result Reproduction

Here is how the model was loaded:

```
model = EnergyRegressor(FEAT_DIM).to(DEVICE)
```

And here is how the trained weights were loaded:

```
model.load_state_dict(torch.load('../trained_model_weights.pth', map_location=DEVICE))
```

Here is the output generated when we run the script:

Using device: cuda

--- Model Architecture ---

```
EnergyRegressor(  
  (net): Sequential(  
    (0): Linear(in_features=784, out_features=4096, bias=True)  
    (1): ReLU(inplace=True)  
    (2): Linear(in_features=4096, out_features=2048, bias=True)  
    (3): ReLU(inplace=True)  
    (4): Linear(in_features=2048, out_features=1024, bias=True)  
    (5): ReLU(inplace=True)  
    (6): Linear(in_features=1024, out_features=512, bias=True)  
    (7): ReLU(inplace=True)  
    (8): Linear(in_features=512, out_features=256, bias=True)  
    (9): ReLU(inplace=True)  
    (10): Linear(in_features=256, out_features=128, bias=True)  
    (11): ReLU(inplace=True)  
    (12): Linear(in_features=128, out_features=64, bias=True)  
    (13): ReLU(inplace=True)  
    (14): Linear(in_features=64, out_features=32, bias=True)  
    (15): ReLU(inplace=True)  
    (16): Linear(in_features=32, out_features=16, bias=True)  
    (17): ReLU(inplace=True)  
    (18): Linear(in_features=16, out_features=8, bias=True)  
    (19): ReLU(inplace=True)  
    (20): Linear(in_features=8, out_features=4, bias=True)  
    (21): ReLU(inplace=True)
```

```
(22): Linear(in_features=4, out_features=2, bias=True)
(23): ReLU(inplace=True)
(24): Linear(in_features=2, out_features=1, bias=True)
)
)
-----

Loading dataset from ../A4_test_data.pt...
Dataset loaded in 0.17s. Shape: x=torch.Size([100000, 784]), energy=torch.Size([100000,
1])

--- Test Results ---
Loss: 288.1554
--- Script Finished ---
```

As shown in the output above, the model and dataset were loaded successfully. The model architecture is a feedforward neural network with 24 layers, and the dataset contains 100,000 samples. The loss value of 288.1554 indicates the performance of the model on the test dataset.

Task 1: MCMC Sampling Implementation

Apart from calculating the acceptance probability, and the burn in time, we calculate the mean probability of the samples which are generated. This is basically the average of $e^{-E(x)}$, where x is the sample (the expression is un-normalized). This lets us know, whether after the burn-in, we were able to reach the high probability regions or not.