

CS726: Programming Assignment 4

April 3, 2025

General Instructions

1. Plagiarism will be strictly penalized including but not limited to reporting to DADAC and zero in assignments. If you use tools such as ChatGPT, Copilot, you must explicitly acknowledge their usage in your report. If you use external sources (e.g., tutorials, papers, or open-source code), you must cite them properly in your report and comments in the code.
2. Submission: Submit a report explaining your approach, implementation details, and results. Clearly mention the contributions of each team member in the report. Submit your code and report as a compressed `<TeamName>_<student1rollno>_<student2rollno>_<student3rollno>.zip` file. Fill a student roll number as NOPE if less than 3 members.
3. Start well ahead of the deadline. Submissions up to two days late will be capped at 80% of the total marks, and no marks will be awarded beyond that.
4. Do not modify the environment provided. Any runtime errors during evaluations will result in zero marks. `README.md` provides instructions and tips to set up the environment and run the code.
5. For most of the assignment, you have to fill in your code in already existing files. Apart from the report, do not submit any additional models and files unless explicitly asked. Any additional files should be placed in the top-level directory.

```
cs726_assgmt4/  
+- TASK-0-1/  
    +- get_results.py  
    +- sampling_algos.py  
+- task-02/  
    +- data/  
    +-causalGraphDiscovery.py  
+- task-03  
    +-gaussianEstimate.py  
+- report.pdf [NEW]
```

6. **STRICTLY FOLLOW THE SUBMISSION GUIDELINES.** Any deviation from these guidelines will result in penalties.

Github link: <https://github.com/Nikita12200/CS726-Assignment-04>

Link to drive: This folder has Trained Regressor model and test data for TASK 0

1 Introduction

Energy-Based Models (EBMs) provide a flexible framework for representing complex probability distributions. Instead of defining a normalized probability density directly, an EBM defines an energy function $E_\theta(x)$, often parameterized by a neural network with parameters θ . The probability distribution is then implicitly defined as:

$$p_\theta(x) = \frac{\exp(-E_\theta(x))}{Z_\theta}$$

where $Z_\theta = \int \exp(-E_\theta(x'))dx'$ is the partition function, which is typically intractable to compute.

While training EBMs presents challenges (often addressed by contrastive methods), sampling from a *trained* EBM, $p(x) \propto \exp(-E(x))$, is also non-trivial. Markov Chain Monte Carlo (MCMC) methods are commonly employed for this task. This assignment focuses on using Langevin dynamics-based MCMC algorithms to draw samples from a probability distribution defined by a pre-trained neural network energy function.

2 Provided Materials

You will be provided with the following files:

- **TASK-0-1/get_results.py**: A Python script defining the neural network architecture used as the energy function. And a Tester class.
- **TASK-0-1/sampling_algos.py**: A file where you are supposed to implement sampling algorithms as explained in Task 1.

Assume the input X to the network is a flattened vector (e.g., $X \in \mathbb{R}^{784}$ for flattened MNIST images).

3 Tasks

3.1 Task 0: Environment Setup and Result Reproduction

1. Set up a Python environment with the necessary libraries (NOTE: You can not use any libraries apart from below mentioned libraries:
 - (a) PyTorch
 - (b) Numpy
 - (c) Matplotlib
2. Load the EnergyRegressor model in **get_results.py**.
3. Load the pre-trained weights from **trained_model_weights.pth** into the network instance.
4. Run the provided **get_results.py** script.
5. In your report, include the output generated by **get_results.py**.

3.2 Task 1: MCMC Sampling Implementation

For this task, consider the loaded, pre-trained neural network $NN(X)$ as the energy function $E(X) = NN(X)$. The target probability distribution is $p(X) \propto \exp(-E(X))$. You will implement three MCMC algorithms to sample from $p(X)$.

1. Implement the Algorithms:

Algorithm 1

```

1: Initialize  $X_0$ 
2: for  $t = 0$  to  $N - 1$  do
3:   Compute gradient  $g_t = \nabla_X E(X_t)$ 
4:   Sample noise  $\xi_t \sim \mathcal{N}(0, I)$ 
5:   Propose  $X' = X_t - \frac{\tau}{2}g_t + \sqrt{\tau}\xi_t$ 
6:   Compute gradient at proposal  $g' = \nabla_X E(X')$ 
7:   Compute acceptance probability:

```

$$\log q(X|X') = -\frac{1}{4\tau} \|X - (X' - \frac{\tau}{2}g')\|^2$$

$$\log q(X'|X_t) = -\frac{1}{4\tau} \|X' - (X_t - \frac{\tau}{2}g_t)\|^2$$

$$\alpha = \min(1, \exp(E(X_t) - E(X') + \log q(X_t|X') - \log q(X'|X_t)))$$

```

8:   Sample  $u \sim \text{Uniform}(0, 1)$ 
9:   if  $u < \alpha$  then
10:      $X_{t+1} \leftarrow X'$ 
11:   else
12:      $X_{t+1} \leftarrow X_t$ 
13:   end if
14: end for
15: return samples  $\{X_t\}_{t > \text{burn-in}}$ 

```

Algorithm 2

```

1: Initialize  $X_0$ 
2: for  $t = 0$  to  $N - 1$  do
3:   Compute gradient  $g_t = \nabla_X E(X_t)$ 
4:   Sample noise  $\xi_t \sim \mathcal{N}(0, I)$ 
5:   Update  $X_{t+1} = X_t - \frac{\tau}{2}g_t + \sqrt{\tau}\xi_t$ 
6: end for
7: return samples  $\{X_t\}_{t > \text{burn-in}}$ 

```

2. **Generate Samples:** Run your implementations to generate a chain of samples for each algorithm.
3. Report sampling Time: Time taken to finish burn-in for each algorithm.
4. **Visualization / Qualitative Results:** Use t-SNE or similar method to plot this high dimensional samples in 2d or 3d space to visualize the distribution for both methods.

3.3 Task 2: Implement the ENCO and PC Algorithms for Observational Data

Causal graphs are directed acyclic graphs (DAGs) that depict causal relationships among variables. Each node in the graph corresponds to a variable, and a directed edge from node A to node B signifies that A exerts a direct causal influence on B . Inferring causal graphs from data is a critical challenge across disciplines such as machine learning, statistics, and the sciences.

Various techniques are available for causal discovery, including constraint-based approaches (e.g., PC algorithm), score-based methods (e.g., NOTEARS), and hybrid methods that utilize both observational

and interventional data (e.g., ENCO). In this assignment, you are required to implement both the modified ENCO algorithm and the PC algorithm, tailored to work exclusively with observational data to uncover causal graphs. You will apply these algorithms to predict causal relationships within the provided dataset.

Two sample input datasets and their corresponding output adjacency matrices are provided in the `data` folder for reference:

- `X1.csv`: A sample observational dataset with 499 rows and 15 columns (variables X_1 to X_{15}).
- `W_true.csv`: The true adjacency matrix for `X1.csv`, a 15×15 matrix with weighted edges.
- `predicted_enco_X1.csv`: Sample ENCO-predicted adjacency matrix for `X1.csv`.
- `predicted_pc_X1.csv`: Sample PC-predicted adjacency matrix for `X1.csv`.

These files can guide you in verifying the correctness of your implementations.

3.3.1 Submission Requirement

1. A Python file named `causalGraphDiscovery.py` containing your implementation.
2. The `causalGraphDiscovery.py` file should include an `ENCO` class and PC algorithm functions with the following components:
 - (a) **Data Loading**: Implement logic to load the observational data from `X1.csv` into a suitable format (e.g., numpy array or PyTorch tensor).
 - (b) **ENCO Algorithm Implementation**:
 - Develop the `ENCO` class with methods to initialize the model, perform distribution fitting, and update graph parameters using observational data only.
 - Implement the `discover_graph` method to train the model and return the predicted adjacency matrix.
 - Include detailed comments explaining the neural network structure, loss function (e.g., Gaussian log-likelihood with sparsity penalty), and parameter updates for γ and θ .
 - Save the predicted ENCO adjacency matrix to `predicted_enco_X1.csv`.
 - (c) **PC Algorithm Implementation**:
 - Implement the `partial_correlation`, `independence_test`, and `pc_algorithm` functions to infer the causal graph using constraint-based methods.
 - Use partial correlation and conditional independence tests to remove edges from an initially fully connected graph, followed by edge orientation rules.
 - Include comments explaining the significance level (α), maximum conditioning set size, and the step-by-step process of the PC algorithm.
 - Save the predicted PC adjacency matrix to `predicted_pc_X1.csv`.
 - (d) **Visualization**: Implement a method to visualize the true and predicted adjacency matrices (e.g., using heatmaps) and save them as images (e.g., `true_graph.png`, `enco_graph.png`, `pc_graph.png`).

3.4 Task 3: Approximating a Black-Box Function Using Gaussian Processes

The objective of this task is to approximate a given black-box function using Gaussian Processes (GPs) without relying on machine learning libraries like `scikit-learn`. Instead, you will implement the Gaussian Process regression algorithm manually using only NumPy for numerical computations and Matplotlib for visualization. You will experiment with different kernels, vary the number of training samples, and analyze how these choices affect the GP's ability to model the function. This exercise will provide insights into the flexibility, limitations, and practical implementation challenges of Gaussian Processes in regression tasks.

Black-Box Function: The Branin-Hoo Function

The black-box function to approximate is the Branin-Hoo function, defined as:

$$f(x_1, x_2) = a \left(x_2 - \frac{5.1}{4\pi^2} x_1^2 + \frac{5}{\pi} x_1 - 6 \right)^2 + 10 \left(1 - \frac{1}{8\pi} \right) \cos(x_1) + 10, \quad (1)$$

where $a = 1$, and the input domain is:

$$x_1 \in [-5, 10], \quad x_2 \in [0, 15].$$

Your goal is to approximate this function using a manually implemented Gaussian Process regression algorithm.

Complete the file `gaussianEstimate.py` as follows:

Task Breakdown

- **Implement the Branin-Hoo Function:** Write a Python function to compute the Branin-Hoo function, which will serve as the black-box function.
- **Generate Training Data:** Sample n_{samples} points uniformly from the input space, experimenting with the following values: $n_{\text{samples}} = 10, 20, 50, 100$.
- **Implement Gaussian Process Regression Manually:**
 - Use `NumPy` to implement the Gaussian Process regression algorithm from scratch, including the computation of the posterior mean and variance.
 - Implement the following kernel types manually:
 - * **Radial Basis Function (RBF) Kernel:** $k(x, x') = \sigma_f^2 \exp\left(-\frac{\|x - x'\|^2}{2\ell^2}\right)$, where ℓ is the length scale and σ_f is the signal variance.
 - * **Matérn Kernel ($\nu = 1.5$):**

$$k(x, x') = \sigma_f^2 \left(1 + \frac{\sqrt{3}\|x - x'\|}{\ell} \right) \exp\left(-\frac{\sqrt{3}\|x - x'\|}{\ell}\right),$$

where ℓ is the length scale and σ_f is the signal variance.

- * **Rational Quadratic Kernel:**

$$k(x, x') = \sigma_f^2 \left(1 + \frac{\|x - x'\|^2}{2\alpha\ell^2} \right)^{-\alpha},$$

where ℓ is the length scale, σ_f is the signal variance, and α is the scale mixture parameter (set $\alpha = 1$).

- Optimize the kernel hyperparameters (ℓ , σ_f , and noise variance σ_n^2) by maximizing the log-marginal likelihood using a grid search approach. (optional for better predictions)
- **Visualize and Compare Results:** Use `Matplotlib` to generate plots comparing the GP's predicted mean and uncertainty (standard deviation) against the true Branin-Hoo function. Create filled contour plots with a colorbar for each visualization.
- **Analyze Model Performance:** Examine how the choice of kernel and the number of training samples influence the GP's approximation ability. Discuss the trade-offs between different kernels and the impact of sample size on prediction accuracy and uncertainty.

Submission Requirements

Submit the following materials:

1. Python code (`gaussianEstimate.py`) implementing all tasks, including the manual GP implementation, kernel functions, hyperparameter optimization, and plotting.

2. Plots for each combination of kernel (RBF, Matérn, Rational Quadratic) and sample size ($n_{\text{samples}} = 10, 20, 50, 100$), showing:
 - (a) The true Branin-Hoo function.
 - (b) The GP predicted mean.
 - (c) The GP predicted uncertainty (standard deviation).
3. A brief report analyzing the results, including:
 - A comparison of the performance of different kernels.
 - The effect of varying the number of training samples on the quality of the approximation.
 - Observations about the uncertainty estimates and their reliability.